

Chapter 7 Linear Discriminant Function

Definition

$$g_i : \mathbf{R}^d \rightarrow \mathbf{R} \quad (1 \leq i \leq c)$$

- Useful way to represent classifiers
- One function per category

Decide ω_i
if $g_i(\mathbf{x}) > g_j(\mathbf{x})$ for all $j \neq i$

- Minimum Risk

$$g_i(\mathbf{x}) = -R(\alpha_i | \mathbf{x}) \quad (1 \leq i \leq c)$$

- Minimum

$$g_i(\mathbf{x}) = P(\omega_i | \mathbf{x}) \quad (1 \leq i \leq c)$$

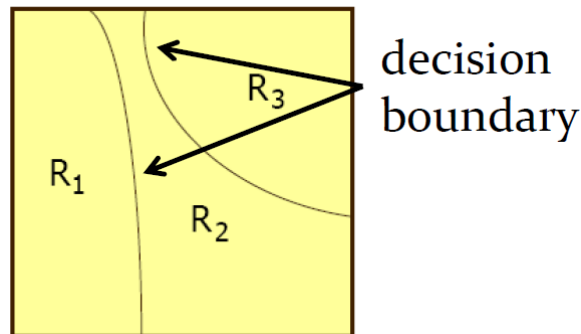
Decision region

$$\mathcal{R}_i = \{\mathbf{x} | \mathbf{x} \in \mathbf{R}^d : g_i(\mathbf{x}) > g_j(\mathbf{x}) \quad \forall j \neq i\}$$

where $\mathcal{R}_i \cap \mathcal{R}_j = \emptyset (i \neq j)$ and $\bigcup_{i=1}^c \mathcal{R}_i = \mathbf{R}^d$

Decision boundary

surface in feature space where ties occur among several largest discriminant functions



Linear Discriminant Functions

$$g_i(\mathbf{x}) = \mathbf{w}_i^t \mathbf{x} + w_{i0}$$

- \mathbf{w}_i : weight vector which is d-dimensional
- w_{i0} : bias/threshold

e.g.

$$\begin{array}{ll} \mathbf{x} = (x_1, x_2, x_3)^t & d = 3, c = 3 \\ g_1(\mathbf{x}) = x_1 - 2x_2 + 4x_3 & \mathbf{w}_1 = (1, -2, 4)^t, w_{10} = 0 \\ g_2(\mathbf{x}) = x_1 + 3x_3 + 4 & \mathbf{w}_2 = (1, 0, 3)^t, w_{20} = 4 \\ g_3(\mathbf{x}) = -2 & \mathbf{w}_3 = (0, 0, 0)^t, w_{30} = -2 \end{array}$$

Generalized Linear Discriminant Function

Quadratic discriminant function (二次判别函数)

$$g(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j$$

Polynomial discriminant function (多项式判别函数)

- e.g. 3rd-order polynomial

$$g(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j + \sum_{i=1}^d \sum_{j=1}^d \sum_{k=1}^d w_{ijk} x_i x_j x_k$$


Nesting version

$$g(\mathbf{x}) = \sum_{i=1}^{\hat{d}} a_i y_i(\mathbf{x}) \Leftrightarrow g(\mathbf{x}) = \mathbf{a}^t \mathbf{y}$$

- \mathbf{a} : the \hat{d} -dimensional weight vector $(a_1, a_2, \dots, a_{\hat{d}})^t$
- \mathbf{y} : the \hat{d} -dimensional transformed feature vector $(y_1(\mathbf{x}), y_2(\mathbf{x}), \dots, y_{\hat{d}}(\mathbf{x}))^t$, where $y_i(\mathbf{x})$ can be viewed as a feature detecting subsystem
- The resulting discriminant function $g(x)$ may not be linear in x , but it is linear in \mathbf{y} .

An example (quadratic in \mathbf{x} \rightarrow linear in \mathbf{y})

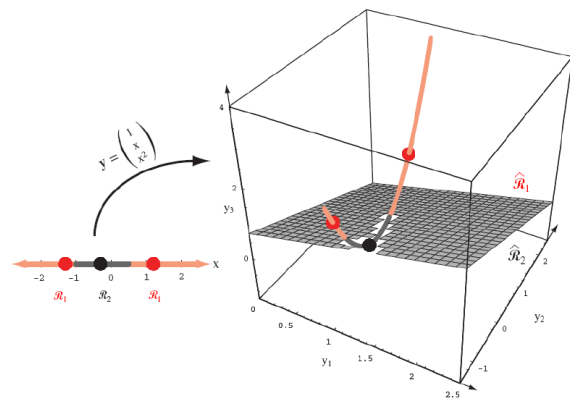
$$g(\mathbf{x}) = a_1 + a_2x + a_3x^2$$



$$\mathbf{a} = (a_1, a_2, a_3)^t$$

$$\mathbf{y} = \begin{pmatrix} y_1(\mathbf{x}) \\ y_2(\mathbf{x}) \\ y_3(\mathbf{x}) \end{pmatrix} = \begin{pmatrix} 1 \\ x \\ x^2 \end{pmatrix}$$

$$g(\mathbf{x}) = \mathbf{a}^t \mathbf{y}$$



$$\mathbf{a} = (-1, 1, 2)^t$$

Augmentation Representation

$$g(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i = w_0 + \mathbf{w}^t \mathbf{x} \Leftrightarrow g(\mathbf{x}) = \mathbf{a}^t \mathbf{y}$$

where the augmented weight vector \mathbf{a}

$$\mathbf{a} = \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{pmatrix} = \begin{pmatrix} w_0 \\ \mathbf{w} \end{pmatrix}$$

the augmented feature vector \mathbf{y}

$$\mathbf{y} = \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{pmatrix} = \begin{pmatrix} 1 \\ \mathbf{x} \end{pmatrix}$$

The Two-Category Case

Training set

$$\mathcal{D} = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\} \quad (\mathbf{y}_i \in \mathbf{R}^{d+1}, \Omega = \{\omega_1, \omega_2\})$$

The task is determine the (augmented) weight vector $\mathbf{a} \in \mathbf{R}^{d+1}$ where $g(\mathbf{x}) = \mathbf{a}^t \mathbf{y}$ can classify all training samples in \mathcal{D} correctly, which is

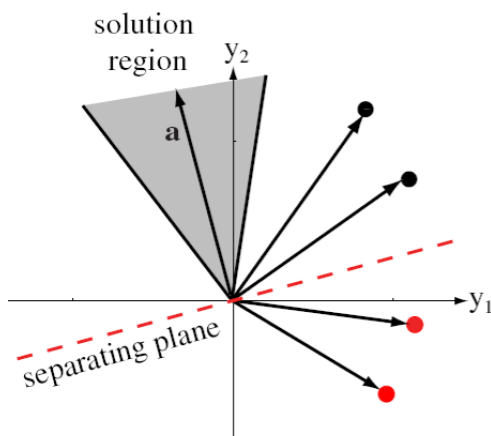
$$\mathbf{a}^t \mathbf{y}_i > 0 \text{ if } \mathbf{y}_i \text{ is labeled } \omega_1$$

$$\mathbf{a}^t \mathbf{y}_i < 0 \text{ if } \mathbf{y}_i \text{ is labeled } \omega_2$$

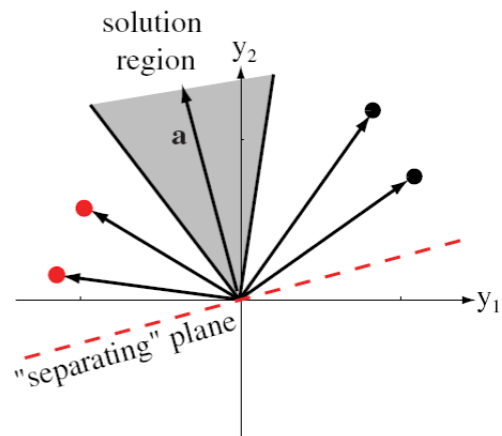
Simplified treatment: "normalization", i.e. replace \mathbf{y}_i with $-\mathbf{y}_i$ if it is labeled ω_2 , that is

$$\mathbf{a}^t \mathbf{y}_i > 0 \text{ for all (normalized) } \mathbf{y}$$

\mathbf{a} should be on the positive side of the hyperplane defined by $\mathbf{a}^t \mathbf{y}_i = 0$ with \mathbf{y}_i being the norm vector



*original
training samples*



*normalized
training samples*

the **Solution Region** is the intersection of n half-spaces yielded by the n training samples,

Gradient Descent

Minimize a criterion/objective function (准则函数) $J(\mathbf{a})$ based on the normalized training samples $\{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\}$ using Gradient Descent

- Basic strategy

$$f(\mathbf{x} + \Delta \mathbf{x}) = f(\mathbf{x}) - \eta \cdot \nabla f(\mathbf{x})^t \cdot \nabla f(\mathbf{x}) + O(\Delta \mathbf{x}^t \cdot \Delta \mathbf{x})$$

$$\mathbf{a}(k+1) = \mathbf{a}(k) - \eta(k) \nabla J(\mathbf{a}(k))$$

1. **begin initialize** **a**, threshold θ , $\eta(\cdot)$, $k \leftarrow 0$
2. **do** $k \leftarrow k + 1$
3. **$\mathbf{a} \leftarrow \mathbf{a} - \eta(k) \nabla J(\mathbf{a})$**
4. **until** $\|\eta(k) \nabla J(\mathbf{a})\| < \theta$
5. **return** **a**
6. **end**

Basic Gradient Descent

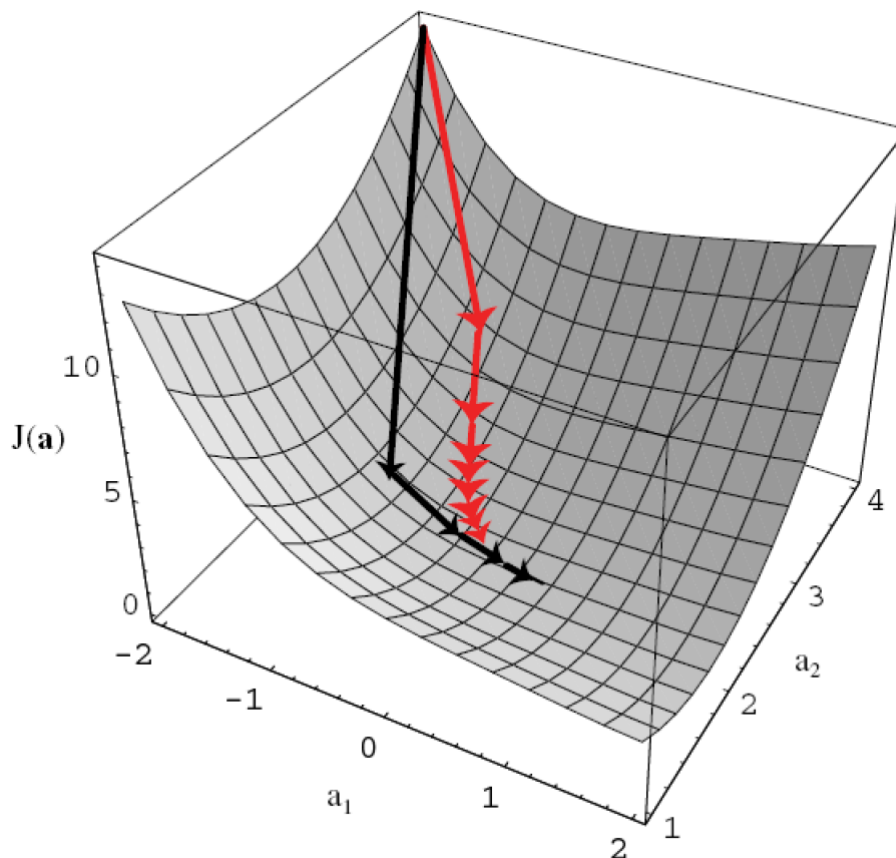
- Newton's Algorithm

$$J(\mathbf{a}) \simeq J(\mathbf{a}(k)) + \nabla J^t(\mathbf{a} - \mathbf{a}(k)) + \frac{1}{2}(\mathbf{a} - \mathbf{a}(k))^t \mathbf{H}(\mathbf{a} - \mathbf{a}(k))$$

$$\mathbf{a}(k+1) = \mathbf{a}(k) - \mathbf{H}^{-1} \nabla J$$

1. **begin initialize** **a**, threshold θ
2. **do**
3. **$\mathbf{a} \leftarrow \mathbf{a} - \mathbf{H}^{-1} \nabla J(\mathbf{a})$**
4. **until** $\|\mathbf{H}^{-1} \nabla J(\mathbf{a})\| < \theta$
5. **return** **a**
6. **end**

- Advantage: better step size than simple gradient descent
- Disadvantage: $O(d^3)$ complexity for matrix inversion; even not applicable if \mathbf{H} is singular

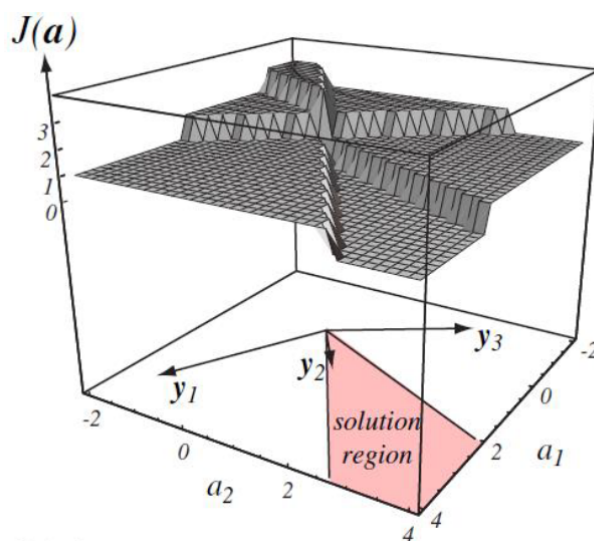


Perceptron Criterion Function

Given the **normalized** training samples $\{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\}$, set the criterion function $J(\mathbf{a})$ as the **number of examples misclassified by \mathbf{a}**

$$J(\mathbf{a}) = \sum_{i=1}^n 1_{\mathbf{a}^t \mathbf{y}_i \leq 0}$$

which is a Piecewise constant function (分段常数函数)



Drawback

- Not compatible with the **gradient descent** procedure for function minimization
 - The gradient is almost 0 (except on the boundary of piecewise region)

A better choice

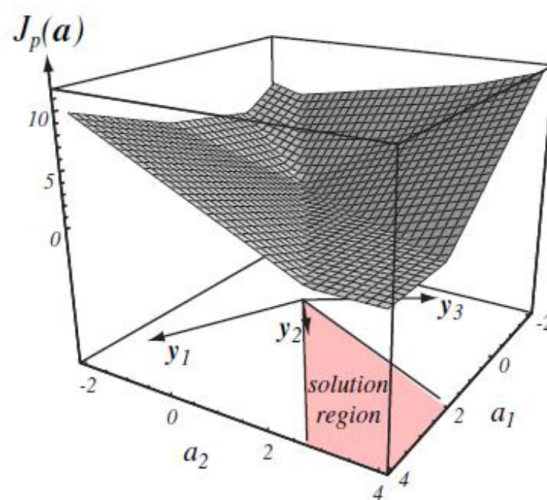
Define

$$\mathcal{Y} = \{\mathbf{y}_i \mid \mathbf{a}^t \mathbf{y}_i \leq 0, 1 \leq i \leq n\}$$

which is the set of samples misclassified by \mathbf{a}

then define criterion function as

$$J_p(\mathbf{a}) = \sum_{\mathbf{y} \in \mathcal{Y}} (-\mathbf{a}^t \mathbf{y})$$



- The gradient

$$\nabla J_p = \sum_{\mathbf{y} \in \mathcal{Y}} (-\mathbf{y})$$

- The iterative update rule

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \sum_{\mathbf{y} \in \mathcal{Y}_k} \mathbf{y}$$

where \mathcal{Y}_k is the set of samples misclassified by $\mathbf{a}(k)$

Batch Perceptron Algorithm

- The next weight vector is obtained by **adding some multiple of the sum of the misclassified samples** to the present weight vector
- Update (correct) \mathbf{a} based on a **set** of misclassified samples

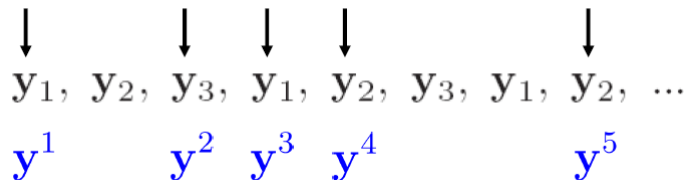
$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \sum_{\mathbf{y} \in \mathcal{Y}_k} \mathbf{y}$$

1. **begin initialize** **a**, threshold θ , $\eta(\cdot)$, $k \leftarrow 0$
2. **do** $k \leftarrow k + 1$
3. $\mathcal{Y}_k = \{\mathbf{y}_i \mid \mathbf{a}^t \mathbf{y}_i \leq 0, 1 \leq i \leq n\}$
4. $\mathbf{a} \leftarrow \mathbf{a} + \eta(k) \sum_{\mathbf{y} \in \mathcal{Y}_k} \mathbf{y}$
5. **until** $\|\eta(k) \sum_{\mathbf{y} \in \mathcal{Y}_k} \mathbf{y}\| < \theta$
6. **return** **a**
7. **end**

Batch Perceptron

Single-Sample Correction

- Update (correct) **a** sequentially whenever it misclassifies a **single sample**
- Consider all training examples **cyclically**
- mark the misclassified samples in **sequence**



Fixed Increment

Setting learning rate $\eta(k) = 1$

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \mathbf{y}^k$$

1. **begin initialize** **a**, $j \leftarrow 0$, $k \leftarrow 0$
2. **do** $j \leftarrow j + 1$
3. $i = ((j - 1) \bmod n) + 1$
4. **if** \mathbf{y}_i is misclassified by **a**
5. **then** $k \leftarrow k + 1$; $\mathbf{y}^k = \mathbf{y}_i$; $\mathbf{a} \leftarrow \mathbf{a} + \mathbf{y}^k$
6. **until** k is kept unchanged for n consecutive rounds
7. **return** **a**
8. **end**

Fixed-Increment Single-Sample Perceptron

We can prove the fact that

- If all training samples are linearly separable
- then the single-sample perceptron converges to a solution vector

Variable Increment

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \mathbf{y}^k$$

1. **begin initialize** \mathbf{a} , margin b , $\eta(\cdot)$, $j \leftarrow 0$, $k \leftarrow 0$
2. **do** $j \leftarrow j + 1$
3. $i = ((j - 1) \bmod n) + 1$
4. **if** $\mathbf{a}^t \mathbf{y}_i \leq b$
5. **then** $k \leftarrow k + 1$; $\mathbf{y}^k = \mathbf{y}_i$; $\mathbf{a} \leftarrow \mathbf{a} + \eta(k) \mathbf{y}^k$
6. **until** k is kept unchanged for n consecutive rounds
7. **return** \mathbf{a}
8. **end**

Variable-Increment
Perceptron with Margin

where we required $\eta(k)$ that

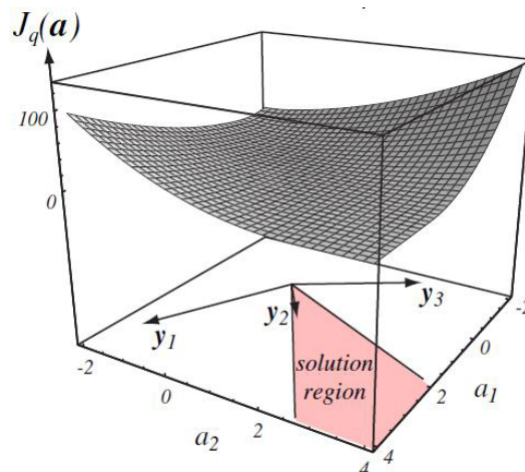
$$\begin{aligned} \eta(k) &\geq 0 \\ \lim_{m \rightarrow \infty} \sum_{k=1}^m \eta(k) &= \infty \\ \lim_{m \rightarrow \infty} \frac{\sum_{k=1}^m \eta^2(k)}{(\sum_{k=1}^m \eta(k))^2} &= 0 \end{aligned}$$

so that it could converge with linearly separable training samples

Other Criterion Functions

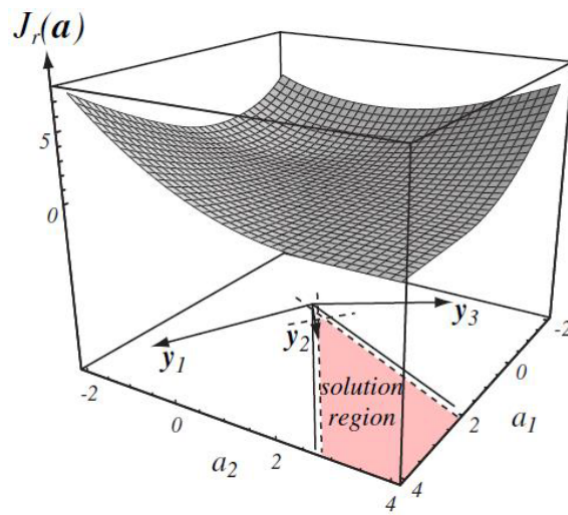
Quadratic Form

$$\begin{aligned} J_q(\mathbf{a}) &= \sum_{\mathbf{y} \in \mathcal{Y}} (\mathbf{a}^t \mathbf{y})^2 \\ \mathcal{Y} &= \{\mathbf{y}_i \mid \mathbf{a}^t \mathbf{y}_i \leq 0, 1 \leq i \leq n\} \\ \nabla J_q &= 2 \sum_{\mathbf{y} \in \mathcal{Y}} (\mathbf{a}^t \mathbf{y}) \mathbf{y} \end{aligned}$$



Distance Form

$$\begin{aligned} J_r(\mathbf{a}) &= \frac{1}{2} \sum_{\mathbf{y} \in \mathcal{Y}} \frac{(\mathbf{a}^t \mathbf{y} - b)^2}{\|\mathbf{y}\|^2} \\ \mathcal{Y} &= \{\mathbf{y}_i \mid \mathbf{a}^t \mathbf{y}_i \leq 0, 1 \leq i \leq n\} \\ \nabla J_r &= \sum_{\mathbf{y} \in \mathcal{Y}} \frac{\mathbf{a}^t \mathbf{y} - b}{\|\mathbf{y}\|^2} \mathbf{y} \end{aligned}$$



Minimum Squared Error (MSE)

$$\mathbf{a}^t \mathbf{y}_i = b_i$$

where

- b_i is some **arbitrarily** chosen **positive** constant
- updating for **all samples**

In other words, we need to solve

$$\mathbf{Y}\mathbf{a} = \mathbf{b}$$

$$\begin{pmatrix} y_{10} & y_{11} & \cdots & y_{1d} \\ y_{20} & y_{21} & \cdots & y_{2d} \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ y_{n0} & y_{n1} & \cdots & y_{nd} \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_d \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}$$

where

- \mathbf{Y} is a $n \times (d + 1)$ matrix
- \mathbf{a} is a $d + 1$ dimensional weight vector
- \mathbf{b} is a n dimensional column vector

Usually $n \geq d + 1$, which means there is no exact solution for \mathbf{a}

Sum-of-squared-error criterion function

$$J_s(\mathbf{a}) = \|\mathbf{Y}\mathbf{a} - \mathbf{b}\|^2 = \sum_{i=1}^n (\mathbf{a}^t \mathbf{y}_i - b_i)^2$$

$$\nabla J_s = \sum_{i=1}^n 2 (\mathbf{a}^t \mathbf{y}_i - b_i) \mathbf{y}_i = 2\mathbf{Y}^t (\mathbf{Y}\mathbf{a} - \mathbf{b})$$

Let $\nabla J_s = 0$, we have

$$\mathbf{a} = (\mathbf{Y}^t \mathbf{Y})^{-1} \mathbf{Y}^t \mathbf{b} = \mathbf{Y}^\dagger \mathbf{b}$$

The Widrow-Hoff rule or LMS (least-mean-squared) rule

For Batch mode

$$\mathbf{a}(k+1) = \mathbf{a}(k) - \eta(k) \mathbf{Y}^t (\mathbf{Y} \mathbf{a}(k) - \mathbf{b})$$

For Single-sample mode

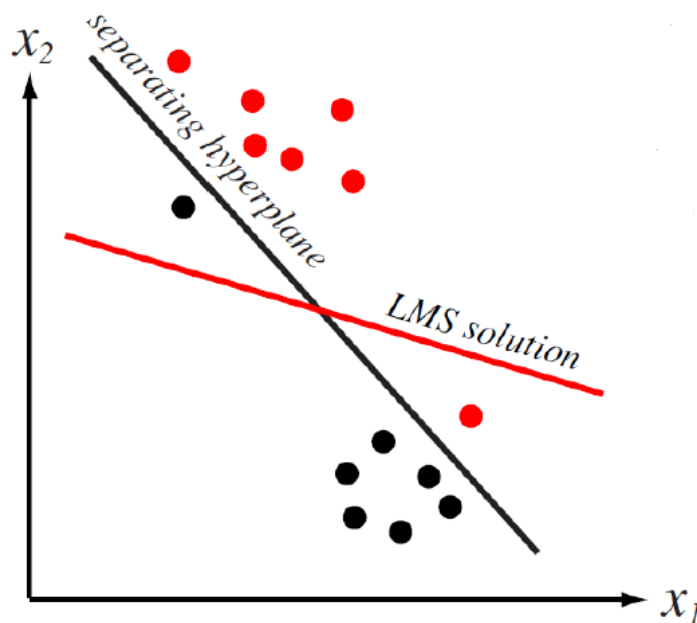
$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) (b_k - \mathbf{a}^t(k) \mathbf{y}^k) \mathbf{y}^k$$

Algorithm

1. **begin initialize** \mathbf{a} , \mathbf{b} , threshold θ , $\eta(\cdot)$, $k \leftarrow 0$
2. **do** $k \leftarrow k + 1$
3. $i = ((k - 1) \bmod n) + 1$
4. $\mathbf{y}^k \leftarrow \mathbf{y}_i$; $b_k \leftarrow b_i$; $\mathbf{a} \leftarrow \mathbf{a} + \eta(k)(b_k - \mathbf{a}^t \mathbf{y}^k) \mathbf{y}^k$
5. **until** $\|\eta(k)(b_k - \mathbf{a}^t \mathbf{y}^k) \mathbf{y}^k\| < \theta$
6. **return** \mathbf{a}
7. **end**

LMS
(least-mean-squared)

- A common choice of $\eta(k)$ for LMS: $\eta(k) = \frac{1}{k}$
- The LMS solution minimizes the sum of squared errors
 - i.e. the (squared) distance of the training samples to the decision hyperplane
 - LMS solution may not be a separating hyperplane



Multi-Category Generalization

$$\mathcal{D} = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\} \quad (\mathbf{y}_k \in \mathbf{R}^{d+1}, \Omega = \{\omega_1, \dots, \omega_c\})$$

Task

Determine the set of (augmented) weight vectors $\{\mathbf{a}_i \in \mathbf{R}^{d+1} \mid 1 \leq i \leq c\}$ where $g_i(\mathbf{x}) = \mathbf{a}_i^t \mathbf{y} (1 \leq i \leq c)$ can classify all training samples in \mathcal{D} correctly

if $\mathbf{y}_k (1 \leq k \leq n)$ is labeled ω_i

$$\forall j \neq i : \mathbf{a}_i^t \mathbf{y}_k > \mathbf{a}_j^t \mathbf{y}_k$$

Keslar's Construction

Transform the **multi-category** classification problem into the **binary** classification problem

